

```
In [1]: # =====
# Notebook setup
# =====

%load_ext autoreload
%autoreload 2

# Control figure size
figsize=(14, 4)

from util import util
from matplotlib import pyplot as plt
import numpy as np
import seaborn as sn
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import GridSearchCV
import xgboost
from sklearn.inspection import permutation_importance

# Generate synthetic data
data, name_map = util.generate_data(size=500, seed=42)
num_cols = [c for c in data.columns[:-1] if len(data[c].unique()) > 2]
cat_cols = [c for c in data.columns[:-1] if len(data[c].unique()) == 2]

# Data pre-processing
X, y = data[data.columns[:-1]].copy(), data[data.columns[-1]].copy()
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])
```

## Non-Linear Models and Importances

### Dealing with Non-Linearities

We'll start by switching to a non-linear model

By doing so:

- We can still account for non-linear correlations
- We can account for interactions among variables
- We might reach a much better accuracy
- ...And hence have a more representative proxy model

Of course there is a price to pay

- Non-linear models are less easy to interpret
- ...And they are at a much higher risk of overfitting

## Gradient Boosted Trees Model

### We'll train a Gradient Boosted Trees model

We'll rely on the Extreme Gradient Boosting package ([XGBoost](#)) for this

```
In [2]: base_est = xgboost.XGBClassifier(tree_method='hist', importance_type='total',
param_grid={'max_depth': [2, 3, 4], 'n_estimators': list(range(20, 41, 5)),
gscv = GridSearchCV(base_est, param_grid=param_grid)
gscv.fit(X, y)
xbm, xbm_params = gscv.best_estimator_, gscv.best_params_
```

**XGBoost is a library for fast, distributed, training of GBT models**

It has support for *multiple loss functions*

- For classification, the default is "reg:logistic", i.e. binary cross-entropy

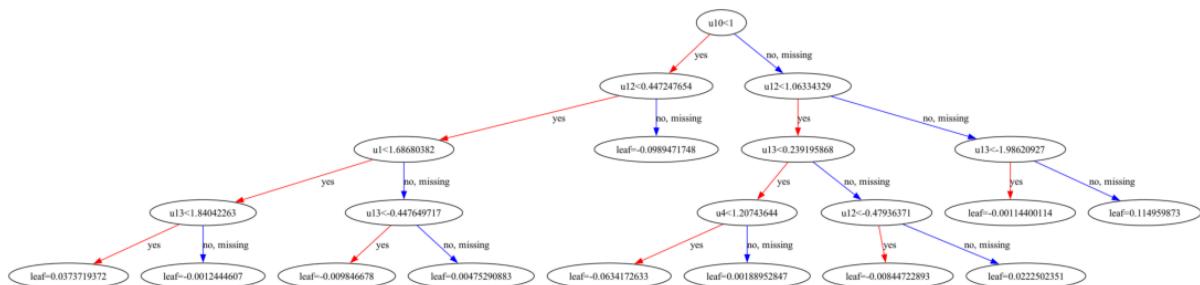
...And for *regularization* (often missing in tree-based models)

- The "reg\_lambda" parameter refers to the weight of an L2 regularization term
- ...Which in GBT is applied **to the leaf labels**

## Gradient Boosted Trees Model

**It's easier to see how regularization work by checking a tree in the ensemble**

```
In [5]: plt.figure(figsize=figsize)
xgboost.plot_tree(xbm, ax=plt.gca(), tree_idx=0);
```



- Assuming  $T$  is the number of leaves and  $w_j$  is the label assigned to each leaf
- ...Then the regularization term is in the form  $\sum_{k=1}^T w_j^2$

## Gradient Boosted Trees Model

## On our dataset, a GBT model has substantially better performance

```
In [6]: xbm_score_cv, xbm_score_test = gscv.best_score_, roc_auc_score(y_test, xbm.predict(X_test))
print(f'AUC score for {xbm_params}: {xbm_score_cv:.2f} (cross-val.), {xbm_score_test:.2f} (test)')
```

AUC score for {'max\_depth': 4, 'n\_estimators': 25, 'reg\_lambda': np.float64(100.0)}: 0.73 (cross-val.), 0.70 (test)

- The AUC score is much higher now
- There is no significant overfitting

**It seems we finally have a model that we can trust**

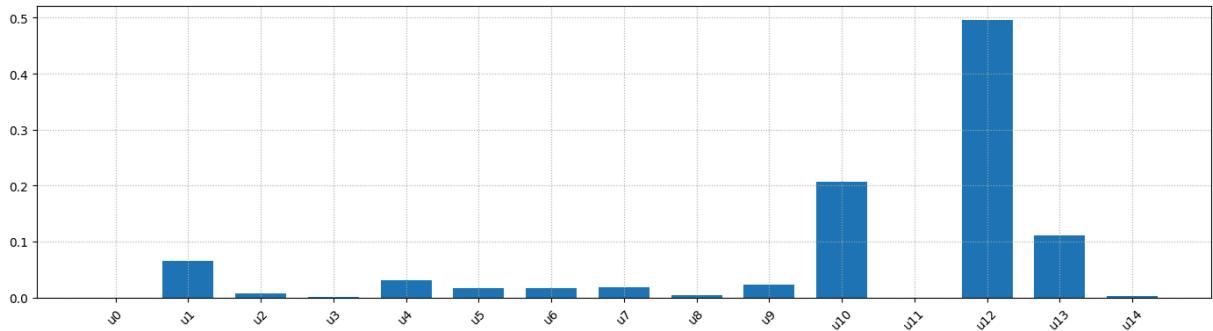
However, we know have an ensemble of many non-linear models

How can we make sense of that?

## Feature Importances

The first option one can probably think of is using *feature importances*

```
In [7]: xbm_imp = pd.Series(index=X.columns, data=xbm.feature_importances_)
util.plot_bars(xbm_imp, figsize=figsize)
```



- The scores differ significantly from those obtained for linear regression (as expected)
- ...But what do they represent?

## Which Feature Importances?

Feature importance is typically presented as this:

- For each input  $x_j$ , we sum the associated *gain* at training time
- Once training is over, we normalize the scores so that they sum up to 1

However, there are other ways to define importance

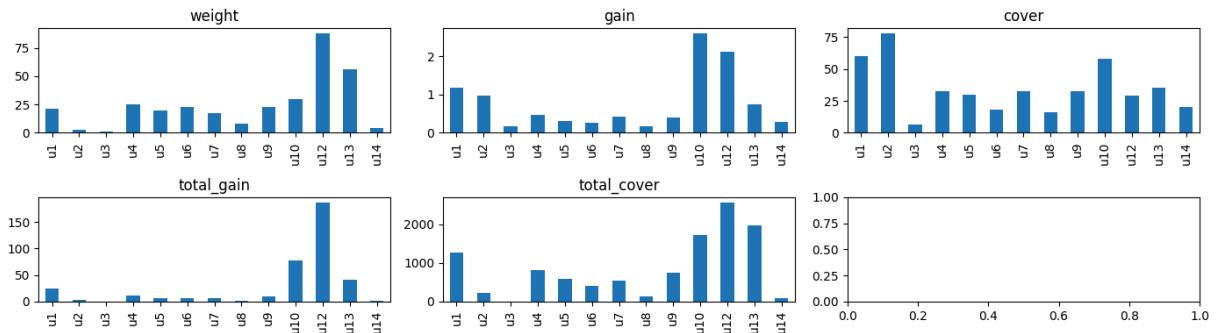
XGBoost supports 5 different approaches:

- "weight": number of times an attribute is used to split
- "gain": average gain associated to splits over an attribute
- "cover": average number of examples for which an attribute is used to decide
- "total\_gain": as above, but replacing the average with a sum
- "total\_cover": as above, but replacing the average with a sum

## Which Feature Importances?

The values of the multiple feature importances can be quite different:

```
In [8]: _, axes = plt.subplots(nrows=2, ncols=3, figsize=figsize)
for ax, imp_type in zip(axes.ravel(), ['weight', 'gain', 'cover', 'total_gain', 'total_cover']):
    pd.Series(xbm.get_booster().get_score(importance_type=imp_type)).plot.bar(ax=ax)
plt.tight_layout()
```



## Importance and Data

Moreover, most importance scores are computed w.r.t. a dataset:

E.g. in XGBoost "gain", "cover", "total\_gain", and "total\_cover"

- For this reason, they are not really properties of the model
- ...But rather of the model and a reference sample

This means that the score semantic depends on the reference sample

By default, importances are computed on the training set

...Which means they are susceptible to overfitting

- The model might split on an attribute because it really is importance
- ...But also due to a [spurious correlation](#)

## Permutation Importance

## We can improve things by changing the way we compute importance

Given a reference sample  $\{x_i, y_i\}_{i=1}^m$

- We can evaluate the performance of our model on the sample
- ...With that of a modified sample where the  $j$ -th input is *made* unimportant

For example, we can achieve that by *permuting* the values of the input

- This will preserve the distribution of the input
- ...But it will break all its correlations

Then, we look at the change in the model performance

- If it is small, the attribute is really unimportant
- Otherwise, the attribute is important

These scores are known as *permutation importances*

## Permutation Importance

Permutation importances are robust w.r.t. spurious correlations

- We just need to repeat the process multiple times
- ...And record means and standard deviations

It's unlikely that we often get a high score by accident

They allow us to choose our reference sample:

On the *training* set, the model might have overfit over the data

- The performance gap will be wider
- ...And the score will reflect how the model *is using* the data

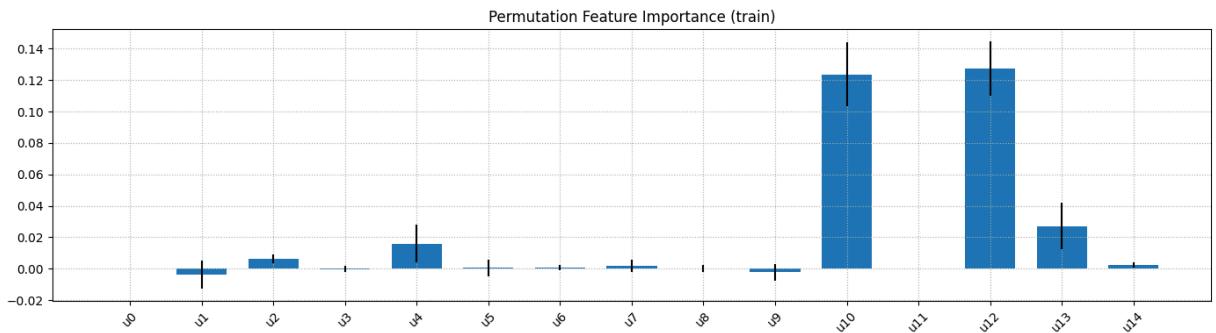
On the *test set*, overfitting will make less of a difference

- ...And the score will reflect how correlated the attribute is with the target

## Permutation Importances, on our Example

Let's check the *training* permutation importances in our case study

```
In [9]: r_train = permutation_importance(xbm, X_train, y_train, n_repeats=30, random_state=42)
xbm_p_imp = pd.Series(index=X.columns, data=r_train.importances_mean)
util.plot_bars(xbm_p_imp, figsize=figsize, std=r_train.importances_std, title="Permutation Importances")
```

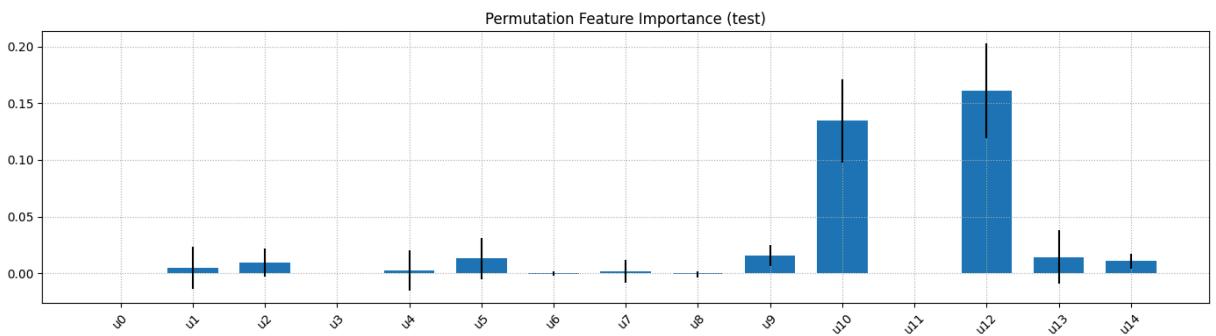


- The closely resemble those XGB "total\_gain", but they are more sparse

## Permutation Importances, on our Example

Let's check the **test** permutation importances in our case study

```
In [10]: r_test = permutation_importance(xbm, X_test, y_test, n_repeats=30, random_state=42)
xbm_p_imp = pd.Series(index=X.columns, data=r_test.importances_mean)
util.plot_bars(xbm_p_imp, figsize=(10, 6), std=r_test.importances_std, title="Permutation Feature Importance (test)")
```



- A few low-importance features become even less relevant on the test data